

# **560.602 GPU/CPU Programming for Engineers**

(<http://www.ce.jhu.edu/dalrymple/classes/602>)

[Prof. Robert A. Dalrymple](#)  
[rad@jhu.edu](mailto:rad@jhu.edu)

201 Latrobe Hall  
410 516 7923

Fall, 2014

**Meeting Time:**

T Th 4:30-5:45

**Location:**

Krieger 307

## **Goal:**

You will learn to write C++ computer programs. You will also learn to write them, using the full capability of multi-core CPUs and Nvidia GPUs for large computational speed-ups. Tools you will use include OpenMP for CPU programming, and CUDA for GPU programs.

## **Grading:**

Homework/Classwork	20%
Midterm Exam	30
Programming Project	25
Final Exam	25

## **Topics**

- C++ Programming
- OpenMP--multi-core parallel programming
- CUDA--GPU Programming
- CUDA Thrust Programming
- Using CUDA libraries

## **Web Resources and References**

- [C++ Language](#)
- [C++ Overview](#)
- [Thinking in C++, 2nd Ed., Vols 1 & 2, Bruce Eckel](#)
- [C++ Library Reference](#)
- [Introduction to Parallel Computing, Blaise Barney](#)
- [POSIX Threads Programming, Blaise Barney](#)
- [Guide to OpenMP: Easy Multithreading Programming for C++, Joel Ylliluoma](#)
- [Intel: Getting Started with OpenMP, Richard Gerber](#)
- [Intel: More Work-Sharing with OpenMP, Richard Gerber](#)
- [Intel: Advanced OpenMP Programming, Richard Gerber](#)
- [GNU libgomp, the GNU OpenMP library](#)
- [An Easy Introduction to CUDA C and C++, Mark Harris](#)
- [Nvidia CUDA Documentation and Guides](#)

# Programming Project: Options

Convert an existing program to parallel

Write a new research program for GPU

Write a new GPU program for me

Project Write up:

- The problem

- The methodology you used


- Optimizations used

- Speed-ups obtained

Project Presentation:

- In-class

- 15 minutes



Ebook: download html  
and code

Start reading->chap 3

<http://mindview.net/Books/TICPP/ThinkingInCPP2e.html>


# Need for parallel programming

Scientific/engineering problems can require immense computation: computation fluid dynamics, climate, meteorology, ocean circulation

Data sets are becoming much larger than computers




# CPU



CPU scaling showing transistor density, power consumption, and efficiency. Chart originally from *The Free Lunch Is Over: A Fundamental Turn Toward Concurrency in Software*

# CUDA C Programming Guide

Theoretical GFLOP/s



## Computer performance

Name	FLOPS
yottaFLOPS	10 <sup>24</sup>
zettaFLOPS	10 <sup>21</sup>
exaFLOPS	10 <sup>18</sup>
petaFLOPS	10 <sup>15</sup>
teraFLOPS	10 <sup>12</sup>
gigaFLOPS	10 <sup>9</sup>
megaFLOPS	10 <sup>6</sup>
kiloFLOPS	10 <sup>3</sup>

# CPU vs GPU

(de acuerdo a Adam y Jamie)



# Flynn's Taxonomy

		Data Stream	
		Single	Multiple
Instruction Stream	Multiple	SISD	SIMD
	Single	MISD	MIMD

Flynn, 1996



# CPU vs Multiprocessor



Centralized Multiprocessor

# Different Approaches to Parallel Programming

## CPU

Threads

OpenMP

OpenMPI

## GPU


CUDA C++, C, Fortran, Python (PyCUDA)

Thrust

Thrust/OpenMP

ACC (new)

# GPUSPH: GPU computing



# Example: Compute Pi


$$\int_0^1 \frac{4}{1+x^2} dx = \pi$$

which is digitized to a summation over N intervals

$$w \sum_{i=0}^{N-1} \left( \frac{4}{1+(i+w/2)^2} \right) \text{ where } w = 1/N$$

# Rectangular (Midpoint) Rule

Divide the interval into N equal pieces,  $h = (b-a)/N$



Evaluate area of rectangle for each piece and sum

$$\int_a^b f(x)dx = h \sum_{i=0, N-1} f\left(a + \left(i + \frac{1}{2}h\right)\right)$$

The bigger N, the smaller the interval and the better answer



```

#include <stdio.h>
#include <stdlib.h> //this library includes atol
#include <time.h>
#include <sys/time.h>
double getTimeElapsed(struct timeval end, struct timeval start)
{
    return (end.tv_sec - start.tv_sec) + (end.tv_usec - start.tv_usec) / 1000000.00;
}

int main (int argc, char *argv[])
{
    double PI25D= 3.141592653589793238462643;
    double x, w, totaltime, pi;
    int i, N;

    // timing info variables
    struct timeval t0, t1;
    double htime;

    pi=0.0;

    N=512;

    if (argc >1) //read in the Number of intervals
    {
        N= atol(argv[1]); //parse command line for number of intervals

        printf("Number of intervals: %d \n",N);
    }
    else
    {
        printf("be sure you had an arg on command line for N greater than 512!\n");
        printf("format: N of intervals \n Default is now 512\n");
    }
}

```

```
w = 1.0f/((double)N); // width of each computational segment
printf(" w = %f \n",w);
```

```
gettimeofday(&t0, NULL);
```

```
for (i=0; i<N ; i++)
{
    x = (i+0.5)*w;
    pi += 4.0/( 1.0 + x*x);
}
```

```
gettimeofday(&t1, NULL);
```

```
hTime=getTimeElapsed(t1,t0);
printf("time for parallel computation: %f \n", hTime);
pi *= w;
printf("pi approx. %1.20f; error = %g \n", pi, PI25D-pi);
}
```

# Examples for Pi: Computation Time

---

<u>Program</u>	<u>N=100</u>	<u>10000</u>	<u>10,000,000</u>
MyPinC	0.00000015	0.000010	0.047
MyPiOpenMP	0.000028	0.000027	0.013
MyPiGPU*	0.0000031	0.000035	0.00004
ThrustPi*	0.000012	0.000033	0.00055

\* GPU library

# Consider No. of Processors

myPiOpenMP 100,000,000 intervals

<u>No. of Processors</u>	<u>Time</u>
2	0.207
4	0.15
6	0.126
8	0.112

Not a linear relationship between # processors and time

# What you need

Computer

Text Editor

C++ compiler with OpenMP

CUDA nvcc compiler  
and video driver



# Mac OS X

Dock: Applications

Finder (on Dock): Directories and Files  
See directory: Applications

TextWrangler (on Dock): Text Editor

Terminal: (on Dock?): in Applications/Utilities

C++ compiler: `/usr/local/bin/g++` `fname`

# Simple UNIX commands

pwd	present working directory
cd name	change directory to name
cd ..	change directory up one level
rm name	remove file named name
ls	list files in directory
ls -l	list files with more information
mv fname fname2	change name
cp fname fname1	copy fname to fname1
cat fname	prints out contents of fname
man unix-command	gives help

# Your first C++ program

Make a directory: CPP

Open TextWranger

```
#include <iostream>
using namespace std;
```

```
int main()
{
    cout << "Hello World";
    return 0;
}
```

Save in CPP as “hello.cpp”

In Terminal: cd CPP

/usr/local/bin/g++ hello.cpp

Type ./a.out