

# HW 4

```
#include <omp.h>
#include <stdlib.h> //atoi
#include <iostream>
#include <iomanip>
#include <fstream>
#include <cmath>
//timing routines
#include <time.h>
#include <sys/time.h>

#define SIZE 256

using namespace std;

//function for elapsed time
double getTimeElapsed(struct timeval end, struct timeval start)
{
    return (end.tv_sec - start.tv_sec) + (end.tv_usec - start.tv_usec) / 1000000.00;
}

int main(int argc, char* argv[])
{
    // timing info variables in main
    struct timeval t0, t1;
    double htime;
    bool print = false;
```

# HW4-2

```
float phi[SIZE][SIZE]={};
float phinew[SIZE][SIZE] = {};
float dx;
int testx = SIZE/2;
int testy = SIZE/2;
int np = 8; //default
float converge;
int i,j;
//set up boundary conditions; y=0
for (i=0; i<SIZE; i++)
{
    phi[i][0] = 1.000000; //all other sides already defined as zero
}
if (argc >1)
{
    np= atoi(argv[1]); //parse command line for number of threads

    printf("Number of threads: %d\n",np);
}
else
{
    printf("be sure you had an arg on command line for np!\n");
    exit(1);
}
// omp_set_dynamic(false);
omp_set_num_threads(np);
```

# HW4-3

```
// set up iterations
gettimeofday(&t0, NULL);
for (int iter= 0; iter<100000; iter++)
{
    cout<<"Iteration: "<<iter<<" ";
    converge=0.0f;
    #pragma omp parallel for private(i) reduction(+:converge)
    for (j=1; j <(SIZE-1); j++)
    {
        for (i= 1; i< (SIZE-1); i++)
        {
            phinew[i][j]=0.25*(phi[i-1][j] + phi[i+1][j] + phi[i][j+1]+phi[i][j-1]);
            converge+=abs(phinew[i][j]-phi[i][j]);
        }
    } //end omp for

    if(iter > SIZE)
    {
        cout <<converge/(SIZE*SIZE)<<endl;
    }

    // update
    #pragma omp parallel for private(i)
    for (j=1; j <(SIZE-1); j++)
    {
        for (i=1; i<(SIZE-1); i++)
            phi[i][j]=phinew[i][j];
    } //end omp for
}
```

```
if (converge/(SIZE*SIZE) < 0.00001)
```

```
{
```

```
    gettimeofday(&t1, NULL);
```

```
    cout<<"Converged: "<<converge/(SIZE*SIZE)<<endl;
```

```
    htime=getTimeElapsed(t1,t0);
```

```
    printf("time for computation: %f \n", htime);
```

```
    ofstream myfile;
```

```
        myfile.open("Laplace.txt");
```

```
    if(myfile.is_open())
```

```
    {
```

```
        cout << "file open for writing\n";
```

```
        for (j=0; j<SIZE; j++)
```

```
            { for (i=0; i<SIZE; i++)
```

```
                {myfile <<phi[i][j] <<" " ;}
```

```
                myfile<<endl;
```

```
            }
```

```
        }
```

```
    else
```

```
        {cout << "file open failure\n";
```

```
        }
```

```
    myfile.close();
```

```
    return iter;
```


```
}
```

```
} //end iteration loop
```

```
    printf("Did not converge\n")
```

```
}
```

# HW4-4



# Using the GPU

Start serial work on CPU

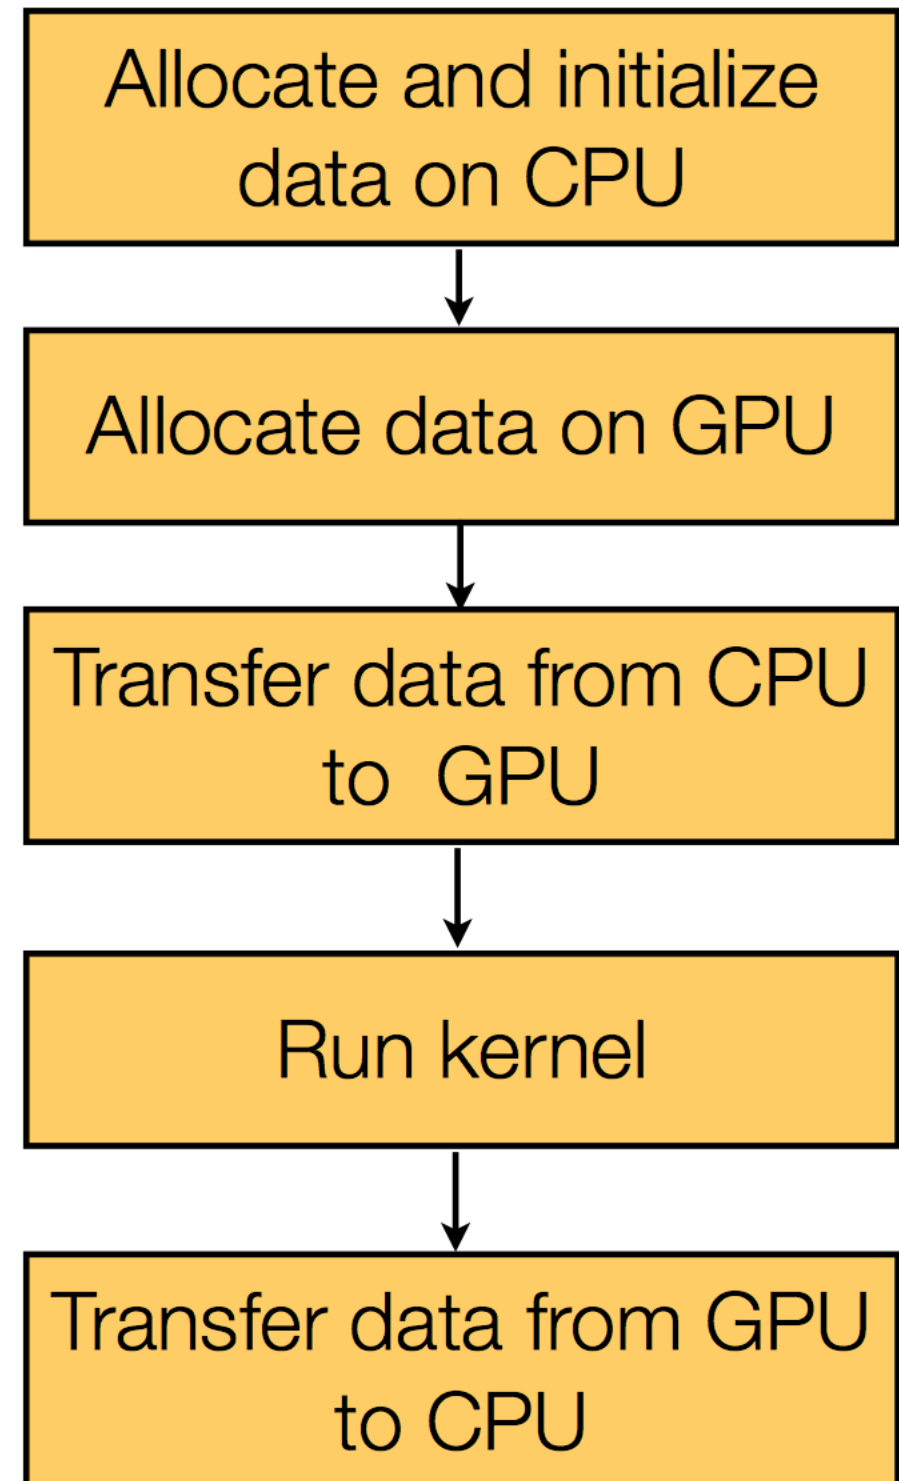
Send data to the GPU

Do computational work on the GPU (kernel)

Pass back answers only when necessary;  
that is, do as much work on GPU as possible

This avoids the PCI-e bus bandwidth issues

Remember: kernels are like functions, but  
massively parallel.



# Data Parallelism

Data parallelism is a program property where many arithmetic operations can be performed on the data simultaneously.

For CUDA programs: part done on CPU; part on the GPU. Parts with no/little data parallelism, do on the CPU; parts with data parallelism, use the GPU.

# HOST (CPU)

## **CUDA routines for:**

memory allocation

data transfer to and from DEVICE

- ordinary data

- constants

- texture arrays (2D read only)

error checking

timing

## C Program Sequential Execution

Serial code

Parallel kernel  
Kernel0<<<>>>()

Serial code

Parallel kernel  
Kernel1<<<>>>()

Host



Device

Grid 0

Block (0, 0)



Block (1, 0)



Block (2, 0)



Block (0, 1)



Block (1, 1)



Block (2, 1)



Host



Device

Grid 1

Block (0, 0)



Block (1, 0)



Block (0, 1)



Block (1, 1)



Block (0, 2)




Block (1, 2)





# Why GPU code can be slower than CPU

Data transfers between devices



# Best Practices

Keep the GPU cores busy. It is often faster to just crunch numbers rather than be clever.

Example: Calculating forces between objects.

For  $N$  objects, calculate the force on object  $(i)$  due to object  $(j)$ . Since force on object  $(j)$  due to object  $(i)$  is equal and opposite, you could neglect half the computations.

But the logic programming/calculations necessary to do this is not worth the time. Compute everything, even if redundant.

# Memory

CPU and GPU have separate memory (DRAM-dynamic random access memory); with the CPU generally having more.

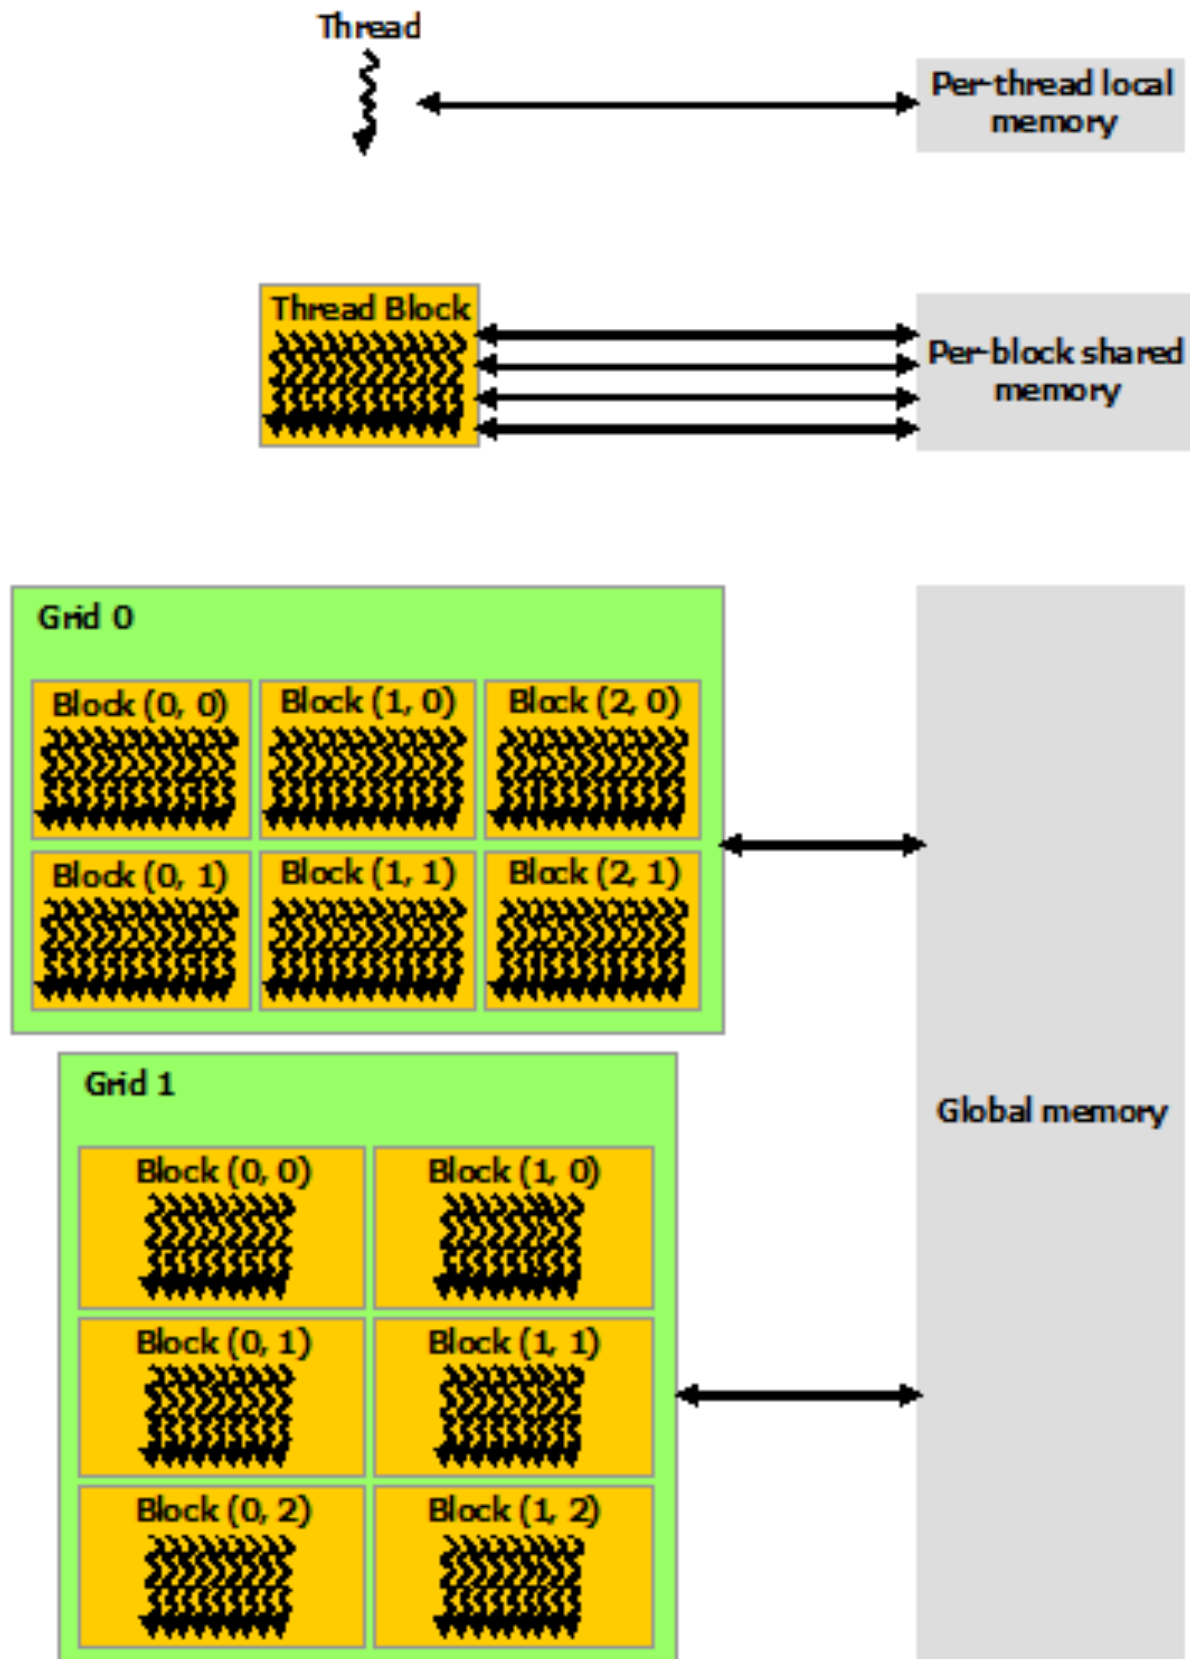
For this reason, doing things on the GPU requires sending data from CPU to GPU, or, in the lingo, Host to Device. And then usually back again.

(Note for compute capable 3.0 devices, can write from GPU. )

Several kinds of Device memory:

Per thread (registers), constant memory, shared (SM), texture, global (slowest).

# Different types of Memory on GPU



# cudaMalloc(&d\_array,num\_bytes)

Allocate (global) memory on device

cudaMalloc needs address of pointer for the allocated memory and the size of the memory required.

Pass-by-pointer (C-style) (see Class3.pdf)

Returns the address of a GPU pointer to the memory

In C,

```
cudaMalloc( (void **) &d_array, num_bytes) //need void ** cast
```

Use `cudaFree(d_array)` to free memory when done

# Error Handling

Every CUDA call returns an error code of type `cudaError_t`

`cudaMalloc( ....)` returns `cudaSuccess`, `cudaErrorInvalidValue`,  
`cudaErrorInvalidDevicePointer`, `cudaInvalidMemcpyDirection`

```
cudaError_t cudaGetLastError(void);
```

Readable error messages can be obtained by

```
char* cudaGetErrorString (cudaError_t error)
```

# Error Checking CUDA calls

/\* CUDA error checking. Usage: after every cuda command, such as kernel launch, type cudaCheckErrors("kernel launch failed") replacing message with the type of cuda command

```
#include <stdio.h>    */
```

```
#define cudaCheckError(msg) \
do { \
    cudaError_t __err = cudaGetLastError(); \
    if (__err != cudaSuccess) { \
        fprintf(stderr, "Fatal error: %s (%s at %s:%d)\n", \
            msg, cudaGetErrorString(__err), \
            __FILE__, __LINE__); \
        fprintf(stderr, "*** FAILED - ABORTING\n"); \
        exit(1); \
    } \
} while (0) // from Robert Crovella, stackoverflow
```

Example:

```
cudaMalloc( (void **) &d_part_sum, N*sizeof(double));
cudaCheckError("cudaMalloc");
```

# Executing kernels

// kernel definition:

```
__global__ void vectorAdd(int *a, int *b, int *c, int n)
{
    int i=threadIdx.x; //replaces for loop as whole function is parallel
    if (i<n)           //only do as many threads as you have data
        c[i] = a[i] + b[i];
}
```

In main:

```
vectorAdd<<<1,SIZE>>>(d_a,d_b,d_c, SIZE); //CUDA kernel call
```

```
kernel <<< num_blocks, threads_per_block>>> (args);
```

so one 1D block of length SIZE, so SIZE threads are used in  
vectorAdd (i = 0, i< SIZE)



# MyPinC —> MyPiGPU

Convert MyPinC to GPU code

```
__device__ __constant__ double PI25D= 3.141592653589793238462643;
```

From DEVICE (to save copying data back):

```
printf("Pi = %.10g, error = %.3g \n", sum*interval_size, interval_size*sum-PI25D);
```

```
__syncthreads();
```

Steps:

1. create a kernel function for integrals
2. create space on device for integrands
3. obtain pi on device
4. print out the answer (from device)

```
nvcc myPiGPU.cu -arch compute_30
```

# Kernel

```
__global__ void term (double* d_part_sum, int n)
{
    double interval_size = 1.0/(double)n;
    int tid = blockIdx.x*blockDim.x + threadIdx.x;
    if (tid < n)
    {
        double x = interval_size * ((double) tid -0.5);
        d_part_sum[tid] = 4.0 / (1.0 + x * x);
    }
    __syncthreads();
    if (tid ==0)
    {
        double sum = 0.0;
        for (int i=0; i<n; i++)
        {
            sum += d_part_sum[i];
        }
        printf("Pi = %.10g, error = %.3g \n", sum*interval_size, interval_size*sum-PI25D);
    }
}
```

# Using nvcc

tcsh

```
setenv PATH /usr/local/cuda/bin:$PATH
```

```
setenv DYLD_LIBRARY_PATH /usr/local/cuda/lib (or location of  
cuda/lib)
```

```
nvcc myPiGPU.cu -o myPiGPU
```

# Compiling CUDA programs

Any program with CUDA programming should be named: **xxxx.cu**

This file can contain both HOST and DEVICE code.

The compiler is **nvcc**, which can handle both C++ and CUDA code. It is capable of compiling all C++ commands for the HOST and most for the DEVICE.

**nvcc xxxx.cu** -> a.out

nvcc xxxx.cu -o xxxxx yields xxxxx as executable file

Compute capability: newer cards have better capabilities

nvcc -code=sm\_30 xxxx.cu means compute\_capability of 3.0

nvcc -m32 xxxx.cu means compile 32bit code, default is 64 bit (-m64)