

Maxi

# Appropriate CUDA

## Hiding Latency

At least 400-600 clock cycles per global memory access

Keep SM busy: 512 per SM

Occupancy: Number of warps running/max warps

Use of `__launch_bounds__(xxx1,xxx2)` per kernel

where xxx1 is max threads per block

xxx2 is min blocks launched at once on SM

This restricts number of registers used so more threads

## Example:

```
__global__ void __launch_bounds__(maxThreadsPerBlock,  
minBlocksPerMultiprocessor) MyKernel(...)  
{ ... }
```

CUDA GPU Occupancy Calculator

# Multiple GPU

# Different GPU cards on same PCIe bus (same workstation)

# Multi-GPU:

Provides more processors to speed-up work  
and/or provides more memory for on-card storage

## Using nvidia-smi (on Linux):

```

++-----+
| NVIDIA-SMI 331.20   Driver Version: 331.20   |
|-----+-----+-----+
| GPU  Name      Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan  Temp  Perf  Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M. |
|=====+=====+=====+
|  0  Quadro K600        Off | 0000:01:00.0   On |          N/A |
| 25%   52C    P0      N/A /  N/A |  211MiB / 1023MiB |    1%      Default |
|-----+-----+-----+
|  1  Tesla K20c         Off | 0000:02:00.0   Off |          0 |
| 30%   39C    P8      15W / 225W |   12MiB / 4799MiB |    0%      Default |
|-----+-----+-----+
|  2  Tesla K20c         Off | 0000:03:00.0   Off |          0 |
| 30%   30C    P8      16W / 225W |   12MiB / 4799MiB |    0%      Default |
|-----+-----+-----+

```

# Multiple GPU

Need inter-GPU communication

Checking device information

```
int deviceCount;
cudaGetDeviceCount(&deviceCount);
int device;
for (device = 0; device < deviceCount; ++device) {
    cudaDeviceProp deviceProp;
    cudaGetDeviceProperties(&deviceProp, device);
    printf("Device %d has compute capability %d.%d.\n",
        device, deviceProp.major, deviceProp.minor);
}
```

# Single CPU thread

All CUDA calls are issued to current GPU

`cudaSetDevice()` sets the current GPU

Example of concurrent execution:

```
cudaSetDevice(0);  
kernel <<<...>>> (....);  
cudaSetDevice(1)  
kernel <<<...>>> (....);
```

Each device has its own stream.

# Using GPU-Specific Streams

```
cudaSetDevice(0);           // Set device 0 as current
cudaStream_t s0;
cudaStreamCreate(&s0);      // Create stream s0 on device 0
MyKernel<<<100, 64, 0, s0>>>(); // Launch kernel on device 0 in s0
cudaSetDevice(1);           // Set device 1 as current
cudaStream_t s1;
cudaStreamCreate(&s1);      // Create stream s1 on device 1
MyKernel<<<100, 64, 0, s1>>>(); // Launch kernel on device 1 in s1
```

// This subsequent kernel launch will fail:

```
MyKernel<<<100, 64, 0, s0>>>(); // Launch kernel on device 1 in s0
```

cudaEventRecord() will fail if the input event and input stream are associated to different devices.

cudaEventElapsedTime() will fail if the two input events are associated to different devices.

# Peer-to-Peer Memory Access

```
cudaSetDevice(0);           // Set device 0 as current
float* p0;
cudaMalloc(&p0, 1024*sizeof(float)); // Allocate memory on device 0
MyKernel<<<1000, 128>>>(p0); // Launch kernel on device 0
cudaSetDevice(1);           // Set device 1 as current
cudaDeviceEnablePeerAccess(0, 0); // Enable peer-to-peer access
                                   // with device 0

// Launch kernel on device 1
// This kernel launch can access memory on device 0 at address p0
MyKernel<<<1000, 128>>>(p0);
```



# Peer-to-Peer Memory Copy

```
cudaSetDevice(0);           // Set device 0 as current
float* p0;
size_t size = 1024 * sizeof(float);
cudaMalloc(&p0, size);       // Allocate memory on device 0
cudaSetDevice(1);           // Set device 1 as current
float* p1;
cudaMalloc(&p1, size);       // Allocate memory on device 1
cudaSetDevice(0);           // Set device 0 as current
MyKernel<<<1000, 128>>>(p0); // Launch kernel on device 0
cudaSetDevice(1);           // Set device 1 as current
cudaMemcpyPeer(p1, 1, p0, 0, size); // Copy p0 to p1
MyKernel<<<1000, 128>>>(p1); // Launch kernel on device 1
```

If `cudaDeviceEnablePeerAccess()` is enabled, host not involved, so faster